

When the Grass Really Is Greener

Three-shot search on an exponentiated Ornstein–Uhlenbeck landscape

Peter Cotton

First version: April 6, 2022. This version: September 3, 2026.

Abstract

A searcher evaluates a rugged landscape — the exponential of an Ornstein–Uhlenbeck path — at three points chosen sequentially, and is paid the value at the final point. The two-evaluation problem has a three-phase closed form: flee somewhere fresh if the first value is below the median, stay put if it is exceptional, and otherwise step a distance set by the mean-reversion scale. The interesting question is the third evaluation: given two locations a bracket apart, should the searcher ever place the final point between them? A rapidity coordinate $\theta = \operatorname{artanh}(e^{-\kappa \Delta t})$, under which bridge means compose additively, reduces the comparison to a hyperbolic identity: an interior bridge point has the same conditional mean as the outside point at the composed rapidity and conditional variance larger by the factor $\cosh(\theta_2 + \theta) / \cosh(\theta_2 - \theta)$. Since the objective rewards variance, the interior point beats its mean-matched outside rival everywhere except at the bracket’s ends. Against the best outside choice the bridge middle wins if and only if $b^2(1 + \rho) - 4b\sqrt{\rho} + 2\rho < 0$, where b is the anchor value and ρ the bracket correlation: revisit the bracket on good news, go forth on bad. Numerically, the optimal three-shot policy widens the bracket relative to the two-shot step in anticipation of the revisit option, which is worth up to about one percent of expected value.

1 The problem: three evaluations of an exponentiated OU path

Let X be a stationary Ornstein–Uhlenbeck process on $\mathbb{R}$ with $\mathbb{E}[X_t] = 0$, unit marginal variance, and covariance kernel

$$k(s, t) = e^{-\kappa|s-t|}. \quad (1)$$

The landscape is $f(t) = \exp(X_t)$: locally continuous, mean-reverting in its logarithm, with occasional high peaks where excursions are magnified by the exponential. A searcher samples f at a first location, which we take to be $t_0 = 0$ by translation invariance, discovering $X_0 = b$. Two further evaluations at locations t_1 and then t_2 follow, each chosen with knowledge of all values so far, and the payoff is the value at the final location,

$$u = \mathbb{E}[f(t_2)] = \mathbb{E}[e^{X_{t_2}}]. \quad (2)$$

Since X_{t_2} is conditionally Gaussian given the observations, with mean μ and variance ν depending on the chosen location,

$$\log \mathbb{E}[e^{X_{t_2}}] = \mu + \frac{1}{2}\nu, \quad (3)$$

and the searcher is paid for conditional mean and conditional variance in equal measure. This is the engine of everything below: exploration has value not only through what it might reveal along the way but directly, because a final location with residual uncertainty carries a lognormal bonus of half its variance.

Two conventions tidy the boundary cases. A searcher may take a location to infinity, which by mean reversion is equivalent to sampling a fresh $N(0, 1)$ value independent of everything seen; and a searcher may revisit an existing location, receiving its known value. Neither changes the substance of any result.

2 Two evaluations: flee, explore, or stay

Given only $X_0 = b$, a candidate second location at distance t has $X_t \sim N(b\lambda, 1 - \lambda^2)$ with $\lambda = e^{-\kappa t} \in (0, 1]$. From (3),

$$\log \mathbb{E}[e^{X_t}] = b\lambda + \frac{1}{2}(1 - \lambda^2) = -\frac{1}{2}(\lambda - b)^2 + \frac{1}{2}(b^2 + 1), \quad (4)$$

maximized over $\lambda \in [0, 1]$ at $\lambda^* = \min\{\max\{b, 0\}, 1\}$. Three phases result.

Proposition 1 (Two-shot policy). The optimal second location and its log-value $\zeta(b)$ are

$$t^* = \begin{cases} \infty & b < 0 \quad (\text{flee}), \\ -\frac{1}{\kappa} \log b & 0 \leq b \leq 1 \quad (\text{explore}), \\ 0 & b > 1 \quad (\text{stay}), \end{cases} \quad \zeta(b) = \begin{cases} \frac{1}{2} & b < 0, \\ \frac{1}{2}(b^2 + 1) & 0 \leq b \leq 1, \\ b & b > 1. \end{cases} \quad (5)$$

A below-median first value is abandoned entirely; an exceptional one is kept; in between, the searcher steps to the distance at which the prior pull toward the anchor exactly matches the anchor's height. The value ζ will serve as the outside benchmark throughout: any strategy that moves beyond all existing observations faces, by the Markov property, a fresh copy of this two-shot problem anchored at the nearest observation.

3 Bridge moments and the rapidity coordinate

Now suppose values $X_0 = b$ and $X_{t_1} = d$ are known, and write $\rho = e^{-\kappa t_1}$ for the bracket correlation. For an interior location $t_2 \in (0, t_1)$ let $\lambda_2 = e^{-\kappa t_2}$ and $\lambda = e^{-\kappa(t_1 - t_2)}$ be the correlations to the two ends, so $\lambda_2 \lambda = \rho$. Gaussian conditioning gives the bridge moments

$$\mu = \frac{\lambda_2(1 - \lambda^2)b + \lambda(1 - \lambda_2^2)d}{1 - \rho^2}, \quad \nu = \frac{(1 - \lambda_2^2)(1 - \lambda^2)}{1 - \rho^2}. \quad (6)$$

The bridge sags: for $b = d > 0$ the mean is below b everywhere strictly inside, pulled toward the process mean, while the variance rises from zero at the ends to a maximum at the middle.

Define rapidities

$$\theta_2 = \text{artanh}(\lambda_2), \quad \theta = \text{artanh}(\lambda), \quad (7)$$

so named because, in the symmetric case $d = b$, bridge means compose like relativistic velocities:

$$\mu = b \frac{\lambda_2 + \lambda}{1 + \lambda_2 \lambda} = b \tanh(\theta_2 + \theta). \quad (8)$$

The variance in (6) also takes a hyperbolic product form. Substituting $1 - \tanh^2 = \operatorname{sech}^2$ and using $\cosh(A + B)\cosh(A - B) = \cosh^2 A \cosh^2 B - \sinh^2 A \sinh^2 B$,

$$\nu = \frac{\operatorname{sech}^2 \theta_2 \operatorname{sech}^2 \theta}{1 - \tanh^2 \theta_2 \tanh^2 \theta} = \operatorname{sech}(\theta_2 + \theta) \operatorname{sech}(\theta_2 - \theta). \quad (9)$$

An outside location beyond t_1 at correlation $\lambda' = \tanh \theta'$ to its anchor has, by the Markov property, mean $d \tanh \theta'$ and variance

$$1 - \tanh^2 \theta' = \operatorname{sech}^2 \theta'. \quad (10)$$

4 Interior points beat their mean-matched outside rivals

Equations (8)–(10) settle the comparison between an interior point and the outside point engineered to match its mean.

Theorem 1 (Bridge dominance). Let $d = b$ and let $t_2 \in (0, t_1)$ be an interior location with rapidities (θ_2, θ) . The outside location at composed rapidity $\theta' = \theta_2 + \theta$ has the same conditional mean $b \tanh(\theta_2 + \theta)$, and the ratio of conditional variances is

$$\frac{\nu_{\text{inside}}}{\nu_{\text{outside}}} = \frac{\operatorname{sech}(\theta_2 + \theta) \operatorname{sech}(\theta_2 - \theta)}{\operatorname{sech}^2(\theta_2 + \theta)} = \frac{\cosh(\theta_2 + \theta)}{\cosh(\theta_2 - \theta)} \geq 1, \quad (11)$$

with equality only in the degenerate limits $\theta_2 \rightarrow \infty$ or $\theta \rightarrow \infty$ (the bracket's ends). By (3) the interior point's utility is at least that of its mean-matched outside rival, strictly so in the interior, for every anchor value b .

Proof. The mean statement is (8) against the outside mean $b \tanh \theta'$ at $\theta' = \theta_2 + \theta$; the variance ratio is (9) against (10); and $\cosh(\theta_2 + \theta) \geq \cosh(\theta_2 - \theta)$ with equality iff $\theta_2 \theta = 0$. Utility (3) is increasing in ν at fixed μ . $\square$

The interior of a bracket whose two ends are both worth b is, point for point, a better bet than the outside locations with the same expected height: the bridge keeps most of the anchors' mean while retaining, at its middle, variance $\operatorname{sech}(\theta_2 + \theta)$ against the outside's $\operatorname{sech}^2(\theta_2 + \theta)$ — up to three times as much residual uncertainty at the same expectation, and the objective pays for uncertainty.

5 The exact boundary: when to revisit, when to go forth

Dominance over mean-matched rivals does not by itself decide the searcher's move, because the outside player is free to choose any correlation, not the mean-matched one. The relevant comparison is the best interior point against the best outside value ζ of Proposition 1. In the symmetric case the middle of the bridge, $\lambda_2 = \lambda = \sqrt{\rho}$, attains the interior optimum, with log-utility

$$U_{\text{mid}}(b, \rho) = \frac{2b\sqrt{\rho}}{1 + \rho} + \frac{1 - \rho}{2(1 + \rho)}. \quad (12)$$




Figure 1: Left: the symmetric-case region (13) where the bridge middle beats every outside choice, in the plane of bracket correlation ρ and anchor value $b = d$. The dashed line is the two-shot-inherited bracket $\rho = b$, which enters the region at $b^* = 0.4233$. Right: the value of the revisit option — the percentage increase in expected payoff of the full three-shot optimum over the best policy forbidden to backtrack — as a function of the first discovered value b , with the bracket optimized in both cases.

Theorem 2 (Revisit boundary). For $d = b \in [0, 1]$ the bridge middle strictly beats every outside location if and only if

$$b^2(1 + \rho) - 4b\sqrt{\rho} + 2\rho < 0, \quad (13)$$

that is, for b strictly between

$$b_{\pm}(\rho) = \frac{\sqrt{\rho} \left(2 \pm \sqrt{2(1 - \rho)} \right)}{1 + \rho}. \quad (14)$$

Proof. For $b \in [0, 1]$ the outside benchmark is $\zeta(b) = \frac{1}{2}(b^2 + 1)$. The inequality $U_{\text{mid}} > \zeta$ becomes, after multiplying by $2(1 + \rho)$ and simplifying, exactly (13); the roots follow from the quadratic formula with discriminant $8\rho(1 - \rho)$. $\square$

The region is nonempty for every $\rho \in (0, 1)$ and, along the two-shot-inherited bracket $\rho = b$, opens at

$$b^* = 0.4233\dots, \quad (15)$$

the root of $U_{\text{mid}}(b, b) = \zeta(b)$: a searcher whose first two evaluations both came in at $b > b^*$ does better returning to the middle of the bracket than striking out anew. Spot values, with the bracket at $\rho = b$: at $b = 0.30$ the middle offers 0.5220 against the outside 0.5450 (go forth); at $b = 0.50$, 0.6381 against 0.6250 (revisit); at $b = 0.90$, 0.9251 against 0.9050 (revisit). By continuity the revisit region extends to $d \neq b$: at $b = 0.5$, $\rho = 0.5$ the interior wins for all $d \geq b$ and loses for $d \leq 0.45$. Good news on both ends of the bracket favors the interior; bad news at either end favors leaving.

b	optimal ρ	value	value, no revisit	revisit worth
-0.5	$\rightarrow 0$	0.8306	0.8306	0.00%
0.1	0.06	0.8398	0.8394	0.03%
0.3	0.16	0.8728	0.8695	0.33%
0.5	0.28	0.9392	0.9304	0.88%
0.7	0.40	1.0378	1.0262	1.16%
0.9	0.50	1.1670	1.1567	1.03%
1.2	0.62	1.4060	1.4004	0.56%

Table 1: The three-shot game: optimal bracket correlation, log-value, the log-value of the best policy forbidden to revisit the bracket interior, and the percentage the revisit option adds to expected payoff. Gauss–Hermite quadrature with 61 nodes over the second value d ; interior optima on a 400-point grid.

6 The three-shot policy, computed

The first location t_1 should anticipate the endgame. Writing $V(b, d, \rho)$ for the best of the outside values $\zeta(b), \zeta(d)$ and the interior optimum, the three-shot value of a bracket ρ is $\log \mathbb{E}_d[e^{V(b, d, \rho)}]$ with $d \sim N(b\rho, 1 - \rho^2)$, evaluated by quadrature. Optimizing over ρ yields the policy. Three features, all visible in Table 1 and Figure 1:

First, a below-median start still means flight: for $b < 0$ the optimal bracket correlation tends to zero (the searcher abandons the anchor, collecting fresh draws) and the revisit option is worthless, since a bridge between an abandoned anchor and whatever follows is a bridge between things not worth bridging.

Second, the optimal bracket is wider than the two-shot step. At $b = 0.5$ the two-shot policy steps to $\rho = 0.5$; the three-shot optimum stretches to $\rho = 0.28$. The wide bracket is bought in anticipation of the endgame: a wider bracket makes the eventual bridge middle more uncertain, and (3) pays half of that uncertainty directly.

Third, the revisit option’s worth peaks near one percent of expected value around $b \approx 0.7$ and vanishes at both extremes — worthless when the start is abandoned, nearly worthless when the start is so good the policy barely moves. Forbidding backtracking is thus a cheap approximation by value, but it is not the optimal policy, and its prescription is qualitatively wrong exactly when both evaluations came in high.

7 Discussion

The proverb warns that the grass only looks greener elsewhere. The analysis says when it truly is. Go forth on bad news: a disappointing anchor is abandoned, a disappointing second value sends the search outward past the better end. Stay between on good news: when both ends of the bracket came in high, the sagging bridge between them retains most of their promise and, at its middle, several times the residual uncertainty of any outside point with the same expectation — and an objective that exponentiates a Gaussian pays for uncertainty at half-weight. The exact frontier between the two regimes is the quadratic (13).

A caution about the word exploration. With one evaluation remaining there is nothing to explore: no discovery at the final point can inform a later choice, because there is none. The

last move is pure placement — put the point where $\mu + \frac{1}{2}\nu$ is largest — and the pull toward uncertain locations is Jensen’s inequality, not curiosity: the objective pays half the variance directly, whether or not anything is ever learned there. Exploration proper enters only with more shots remaining, where early placements buy information. The natural conjecture from the corrected endgame is a two-phase k -shot policy: step outward while values keep climbing, each new high re-anchoring the next step, then refine inside once a high region is bracketed. Settling it is a dynamic program over the posterior, cheap per step with the bridge formulas above.

The rapidity coordinate is the tool worth keeping: bridge means compose additively in $\theta = \text{artanh}(e^{-\kappa\Delta t})$, bridge variances factor as $\text{sech}(\theta_2 + \theta) \text{sech}(\theta_2 - \theta)$, and the inside–outside comparison collapses to the monotonicity of $\cosh$. The same coordinate is the natural discretization for lattice computations of extremes of OU paths, where uniform-in- θ spacing makes correlations additive across cells.

Beyond three shots the closed forms give out, and the k -shot policy is a dynamic program over the posterior — bridge moments per interval, maximization over intervals — that the formulas here make cheap per step. The landscapes themselves remain a useful synthetic test bed for few-shot exploration algorithms: the choice between striking out and revisiting is a one-bit probe of whether a search strategy has grasped the trade between mean reversion and the variance bonus, and (13) is the answer key.

Reproducibility

All numbers and the figure are produced by `numerics.py` alongside this source, with an independent verification of Theorems 1 and 2 by exact conditioning and weighted Monte Carlo in `verify/check_inside.py`, in the [browniansearch](#) repository.

References

- [1] S. Callander. Searching for good policies. *American Political Science Review*, 105(4):643–662, 2011.
- [2] O. Baumann, J. Schmidt, and N. Stieglitz. Effective search in rugged performance landscapes: a review and outlook. *Journal of Management*, 45(1):285–318, 2019.
- [3] J.-B. Grill, M. Valko, and R. Munos. Optimistic optimization of a Brownian. In *Advances in Neural Information Processing Systems 31*, 2018.
- [4] J. M. Calvin. Average performance of a class of adaptive algorithms for global optimization. *Annals of Applied Probability*, 7(3):711–730, 1997.
- [5] R. Agrawal. The continuum-armed bandit problem. *SIAM Journal on Control and Optimization*, 33(6):1926–1951, 1995.