

# When the Grass Is Greener

## Three-Shot Search on an Exponentiated Ornstein–Uhlenbeck Landscape

Peter Cotton

First version: April 6, 2022. This version: September 3, 2026.

### Abstract

A searcher evaluates a random landscape at three points, one at a time, and is paid the value at the final point. The landscape is the exponential of a stationary Ornstein–Uhlenbeck path. We solve the problem exactly. With two evaluations the optimal rule has three regimes: abandon a below-average start, keep an exceptional one, and otherwise step a distance given in closed form. With three evaluations the question is where to commit after observing values  $b$  and  $d$  at bracket correlation  $\rho$ . For  $b = d$ , every interior point of the bracket has the same conditional mean as some exterior point, and conditional variance larger by the constant factor  $(1 + \rho)/(1 - \rho) = e^{2\Theta}$ . The interior problem reduces to a parabola on an interval. The best interior point beats every alternative exactly when  $b$  lies between a quadratic root and  $e^{2\Theta}$ . The optimal interior point is the midpoint for moderate  $b$ ; it splits into a pair that moves toward the ends as  $b$  grows, and reaches them at  $b = e^{2\Theta}$ , above which the searcher should stay at an observed point. The interior option raises expected value by about one percent at most. We verify the landscape assumption on 77,040 WiFi measurements and interpret the three evaluations as an incumbent strategy, one pivot, and a commitment.

## 1 Introduction

How many pivots does a firm get? In randomized trials, about three quarters of startups did not pivot at all over a year, about a fifth pivoted once, and fewer than one in ten pivoted more than once (Camuffo et al. 2020, 2024). Training founders to search more scientifically raised the frequency of a single pivot and lowered the frequency of repeated pivots. In practice, strategic search involves very few evaluations.

Callander (2011) modeled this kind of search as trial and error on an unknown function drawn from a Brownian motion. Each trial is implemented and its outcome observed. Most work in this tradition assumes myopic agents. Bardhi and Callander (2026) list farsighted search with a free choice of location as an open problem.

We solve the smallest such problem exactly. The landscape is the exponential of a stationary Ornstein–Uhlenbeck path. Mean reversion means that performance far from known territory regresses to the industry average. The exponential means that small differences in fit produce large differences in sales, so an untried strategy has extra value on average.

The searcher knows the performance of an incumbent strategy. She may try one more. Then she must commit. With two evaluations the optimal rule is a threshold rule: abandon a below-average incumbent, keep an exceptional one, and otherwise step a distance that we

give in closed form. With three evaluations the question is where to commit after the trial: between the two known strategies, beyond the better one, or at one of them. We compute the answer exactly. Commitment between the two is optimal when both values are good but not exceptional. The commitment point is the midpoint for moderate values and moves toward the better-known end as the values improve. Above an exact ceiling the searcher should stay where she is.

The interior option is worth about one percent of expected value at the optimum. Most of the value of the final pivot comes from the information it produces. On landscapes that decorrelate quickly, it is better to probe far away, and the formulas say so: the interior premium is first order in the bracket correlation.

## 2 Related work

Callander (2008) introduced the Brownian policy landscape. In that paper, delegation to a biased expert survives only when the bias is below  $\sigma^2/2\mu$ . The 2011 papers use the landscape for search. An agent experiments when the incumbent outcome is worse than  $\alpha = \sigma^2/2|\mu|$  and stops otherwise. Search ends almost surely, often at a mediocre policy. Simulated runs change policy between one and four times, which is consistent with the pivot counts above. Later papers derive similar thresholds for precedent (Callander and Clark 2017), agenda control (Callander and McCarty 2024), and risk-averse learning (Callander and Matouschek 2019). In all of these the number of evaluations is endogenous and payoffs accrue every period. In our problem the number of evaluations is fixed at three, only the final placement is paid, and the landscape mean-reverts.

Callander and Hummel (2014) analyze a two-period game in which each of two policymakers can experiment once. The equilibrium has two exact cutoffs, and the first mover sometimes experiments although the status quo already delivers his ideal outcome. The players are strategic and payoffs accrue in both periods. Ganz (2020) studies an organization that allows a manager one offline trial before an online choice. That is a two-evaluation problem in which only the second evaluation is paid, and some of its boundaries are computed numerically. We solve the three-evaluation single-agent problem in closed form.

Other papers obtain farsighted results under other restrictions. Urgun and Yariv (2025) require search to be contiguous; the searcher chooses speed and stops at a closed-form draw-down. Cetemen, Urgun and Yariv (2023) extend this to teams and Wong (2025) to flow payoffs. Bardhi (2024) chooses  $k$  sample locations at once to estimate a project. Garfagnini and Strulovici (2016), Carnehl and Schneider (2025), and Callander, Lambert and Matouschek (forthcoming) study infinite sequences of short-lived agents. Aybas and Callander (2025) allow Ito diffusion landscapes, including the Ornstein–Uhlenbeck case, in a one-shot cheap-talk game. That is the only mean-reverting landscape we know of in this literature. Bardhi (2024) also proves that the OU kernel is the only powered-exponential covariance with the Markov property, which supports our choice of landscape.

In the optimization literature the same object appears with large budgets. Calvin (1997) derives average-case convergence rates under Wiener measure. Grill, Valko and Munos (2018) approximate the maximum of a Brownian path. Grosse, Zhang and Hennig (2023) extend this to Gaussian-process samples and prove regret rates. Alabert and Caballero (2018) compute the law of the minimum of a conditioned bridge in order to compare fixed sampling designs. Malladi (2025) and Banchio and Malladi (2025) solve worst-case versions with no prior and a

Lipschitz bound in place of volatility. Hodgson and Lewis (2025) estimate a Gaussian-process search model on clickstream data and solve the policy numerically. We are not aware of a closed-form solution for a small fixed budget of adaptive evaluations under a Bayesian prior.

Glick and Myers (2015) tested the one-evaluation problem on lab subjects. Subjects respond to the Brownian structure but modify too much when the variance is high.

### 3 The model: three evaluations of an exponentiated OU path

Let  $X$  be a stationary Ornstein–Uhlenbeck process on  $\mathbb{R}$  with  $\mathbb{E}[X_t] = 0$ , unit marginal variance, and covariance kernel

$$k(s, t) = e^{-\kappa|s-t|}. \quad (1)$$

Exponential decay of correlation with distance is a standard empirical model, not an assumption special to this paper: it has described shadow fading in radio engineering since Gudmundson (1991). On 77,040 WiFi measurements across a university floor (Feng, Nguyen and Luo 2024), we find the detrended log signal close to Gaussian with spatial correlation fitting an exponential at  $R^2 = 0.98$ ; the scripts are in the paper’s repository. The same shape is standard in several other fields. In Gaussian-process modeling the exponential kernel is the Matérn class at  $\nu = 1/2$  and is the covariance of the OU process (Rasmussen and Williams 2006, §4.2.1). Practical Bayesian optimization prefers Matérn kernels because squared-exponential sample paths are too smooth for real objectives (Snoek, Larochelle and Adams 2012). Weinberger (1990) measures the ruggedness of a fitness landscape by the exponential decay rate of a random walk’s autocorrelation and shows that Kauffman’s NK model is generically of this type. In evolutionary biology the OU process is the standard model of a trait held near an adaptive optimum (Hansen 1997; Butler and King 2004). The landscape is  $f(t) = \exp(X_t)$ : locally continuous, mean-reverting in its logarithm, with occasional high peaks where excursions are magnified by the exponential. A searcher samples  $f$  at a first location, which we take to be  $t_0 = 0$  by translation invariance, discovering  $X_0 = b$ . Two further evaluations at locations  $t_1$  and then  $t_2$  follow, each chosen with knowledge of all values so far, and the payoff is the value at the final location,  $u = \mathbb{E}[f(t_2)]$ . Conditional on the observations gathered by the time  $t_2$  is chosen,  $X_{t_2}$  is Gaussian with some mean  $\mu$  and variance  $\nu$ , and

$$\log \mathbb{E}[e^{X_{t_2}} \mid \mathcal{F}] = \mu + \frac{1}{2}\nu, \quad (2)$$

where  $\mathcal{F}$  denotes the observed locations and values. (The unconditional distribution after adaptive placement need not be Gaussian; every use of (2) below is conditional.) The searcher is thus paid for conditional mean and conditional variance in equal measure: a final location with residual uncertainty carries a lognormal bonus of half its variance.

Figure 1 shows the decision this paper solves. Two conventions tidy the boundary cases. A searcher may take a location to infinity, which by mean reversion is equivalent to sampling a fresh  $N(0, 1)$  value independent of everything seen; and a searcher may finish at an observed location, receiving its known value. Neither changes the substance of any result.




Figure 1: The problem. Two locations have been evaluated and returned equal values (dots); the shaded region is the posterior band for the landscape and the grey curve one landscape consistent with the data. The final evaluation is the consequential one: settle between the anchors, step out past an end, or stay at an observed point. The paper computes which, exactly, as a function of the anchor value and the bracket correlation.

## 4 Two evaluations: flee, explore, or stay

Given only  $X_0 = b$ , a candidate second location at distance  $t$  has  $X_t \sim N(b\lambda, 1 - \lambda^2)$  with  $\lambda = e^{-\kappa t} \in (0, 1]$ . From (2),

$$\log \mathbb{E}[e^{X_t} | X_0 = b] = b\lambda + \frac{1}{2}(1 - \lambda^2) = -\frac{1}{2}(\lambda - b)^2 + \frac{1}{2}(b^2 + 1), \quad (3)$$

maximized over  $\lambda \in [0, 1]$  at  $\lambda^* = \min\{\max\{b, 0\}, 1\}$ .

*Proposition 1* (Two-shot policy). The optimal second location and its log-value  $\zeta(b)$  are

$$t^* = \begin{cases} \infty & b < 0 \quad (\text{flee}), \\ -\frac{1}{\kappa} \log b & 0 \leq b \leq 1 \quad (\text{explore}), \\ 0 & b > 1 \quad (\text{stay}), \end{cases} \quad \zeta(b) = \begin{cases} \frac{1}{2} & b < 0, \\ \frac{1}{2}(b^2 + 1) & 0 \leq b \leq 1, \\ b & b > 1. \end{cases} \quad (4)$$

A below-median first value is abandoned, and an exceptional one is kept: already at two shots, an observed location can be the optimal final location. Practitioners know this fork as pivot or persevere (Ries 2011); the proposition supplies its thresholds and the size of the step when pivoting wins. In between, the searcher steps to the distance at which the prior pull toward the mean exactly balances the anchor's height. The value  $\zeta$  serves as the exterior benchmark throughout: by the Markov property, any placement beyond all existing observations faces a fresh copy of this problem anchored at the nearest observation.

## 5 Bridge moments, and what composes how

Suppose values  $X_0 = b$  and  $X_{t_1} = d$  are known, with bracket correlation  $\rho = e^{-\kappa t_1}$  and rapidity  $\Theta = \text{artanh } \rho$ . For an interior location let  $\lambda_2 = e^{-\kappa t_2}$  and  $\lambda = e^{-\kappa(t_1 - t_2)}$  be the correlations to

the two ends,  $\lambda_2\lambda = \rho$ , with rapidities  $\theta_2 = \text{artanh } \lambda_2$  and  $\theta = \text{artanh } \lambda$ . Gaussian conditioning gives the bridge moments

$$\mu = \frac{\lambda_2(1 - \lambda^2)b + \lambda(1 - \lambda_2^2)d}{1 - \rho^2}, \quad \nu = \frac{(1 - \lambda_2^2)(1 - \lambda^2)}{1 - \rho^2}. \quad (5)$$

For  $b = d > 0$  the bridge sags: the mean is below  $b$  everywhere strictly inside, pulled toward the process mean, while the variance rises from zero at the ends to a maximum at the middle.

A caution on what is additive. Correlations across consecutive intervals multiply,  $\rho(t+s) = \rho(t)\rho(s)$ , so the additive coordinate for composing intervals is  $\kappa t = -\log \rho$ , not the rapidity. The rapidity earns its keep through a different identity: for equal anchors  $d = b$  the bridge mean is

$$\mu = b \frac{\lambda_2 + \lambda}{1 + \lambda_2\lambda} = b \tanh(\theta_2 + \theta), \quad (6)$$

which matches the mean of the exterior point at correlation  $\tanh(\theta_2 + \theta)$  to its anchor. This is a mean-matching device we use next, not a composition law. The bridge variance also takes a hyperbolic product form: substituting  $1 - \tanh^2 = \text{sech}^2$  into (5),

$$\nu = \text{sech}(\theta_2 + \theta) \text{sech}(\theta_2 - \theta), \quad (7)$$

while an exterior point at correlation  $\lambda' = \tanh \theta'$  to its anchor has variance  $1 - \tanh^2 \theta' = \text{sech}^2 \theta'$ .

## 6 Matched-mean dominance at a constant variance ratio

*Theorem 1* (Bridge dominance, equal anchors). Let  $d = b$  and let  $t_2 \in (0, t_1)$  be an interior location with rapidities  $(\theta_2, \theta)$ . The exterior location at composed rapidity  $\theta' = \theta_2 + \theta$  has the same conditional mean  $b \tanh(\theta_2 + \theta)$ , and the ratio of conditional variances is the same at every interior point:

$$\frac{\nu_{\text{inside}}}{\nu_{\text{outside}}} = \frac{\cosh(\theta_2 + \theta)}{\cosh(\theta_2 - \theta)} = \frac{1 + \tanh \theta_2 \tanh \theta}{1 - \tanh \theta_2 \tanh \theta} = \frac{1 + \rho}{1 - \rho} = e^{2\Theta} > 1. \quad (8)$$

Consequently the interior point's conditional distribution dominates its mean-matched exterior rival in the Gaussian convex order, and its expected payoff is at least as large for every convex payoff with finite expectation, strictly so for the exponential.

*Proof.* The mean statement is (6). The ratio of (7) to  $\text{sech}^2(\theta_2 + \theta)$  is  $\cosh(\theta_2 + \theta)/\cosh(\theta_2 - \theta)$ , and expanding both via  $\cosh(A \pm B) = \cosh A \cosh B \pm \sinh A \sinh B$  gives  $(1 + \tanh \theta_2 \tanh \theta)/(1 - \tanh \theta_2 \tanh \theta)$ , which is  $(1 + \rho)/(1 - \rho)$  because  $\tanh \theta_2 \tanh \theta = \lambda_2\lambda = \rho$  on the bracket. Two Gaussians with equal means and ordered variances are ordered in the convex order (the wider equals the narrower plus independent noise), and (2) is the exponential case.  $\square$

The amplification is a property of the bracket, not the position: threefold at  $\rho = \frac{1}{2}$ , nineteenfold at  $\rho = 0.9$ , unbounded as  $\rho \uparrow 1$ . At the bracket's ends both variances vanish; their ratio still tends to  $e^{2\Theta}$ . Dominance over mean-matched exterior rivals, which the theorem explicitly constructs, decides nothing by itself: the exterior player may choose any correlation, and an observed anchor may simply be kept. The object that can fail to exist is different:

an *unvisited* point whose conditional mean matches the best *observed* value. Mean reversion caps unvisited conditional means strictly below a high anchor, and whether the variance bonus makes up the shortfall is precisely the question the phase diagram answers. The next section makes this exact.

## 7 The constrained parabola and the complete symmetric solution

Parametrize the interior by  $x = \lambda_2 \in (\rho, 1)$  and set

$$C = \frac{1+\rho}{1-\rho} = e^{2\Theta}, \quad m = \frac{x+\rho/x}{1+\rho} \in \left[ \frac{2\sqrt{\rho}}{1+\rho}, 1 \right), \quad (9)$$

under which mirror points  $x \leftrightarrow \rho/x$  coincide, the lower endpoint of  $m$  is the bracket middle, and, using  $(1-x^2)(1-\rho^2/x^2) = (1+\rho)^2(1-m^2)$ , the equal-anchor interior log-utility becomes a constrained parabola:

$$L_{\text{inside}}(m) = bm + \frac{C}{2}(1-m^2), \quad m^* = \text{clip}\left(\frac{b}{C}; \frac{2\sqrt{\rho}}{1+\rho}, 1\right). \quad (10)$$

This is the organizing object of the paper: the entire symmetric analysis reads off the position of the unconstrained peak  $b/C$  relative to the interval.

*Theorem 2* (Pitchfork). The bracket middle is the interior optimum iff  $b/C \leq 2\sqrt{\rho}/(1+\rho)$ , i.e.

$$b \leq \frac{2\sqrt{\rho}}{1-\rho} = \sinh 2\tau, \quad \tau = \text{artanh } \sqrt{\rho}. \quad (11)$$

For  $\sinh 2\tau < b < e^{2\Theta}$  the peak is interior and attained by the mirror pair

$$x_{\pm} = \frac{b(1-\rho) \pm \sqrt{b^2(1-\rho)^2 - 4\rho}}{2}, \quad x_+x_- = \rho, \quad (12)$$

with optimal value

$$U_{\text{off}} = \frac{C}{2} + \frac{b^2}{2C} = \frac{1}{2}(b^2e^{-2\Theta} + e^{2\Theta}). \quad (13)$$

For  $b \geq e^{2\Theta}$  the peak  $b/C \geq 1$  lies at or beyond the interval's open end:  $L_{\text{inside}}$  is increasing in  $m$ , no interior maximum is attained, and the interior supremum is the end limit  $b$ .

*Proof.*  $dL/dm = b - Cm$  vanishes at  $m = b/C$ ; the interval endpoints translate via (9) into the stated  $b$ -thresholds ( $b/C = 2\sqrt{\rho}/(1+\rho)$  is  $b = 2\sqrt{\rho}/(1-\rho)$ , and  $b/C = 1$  is  $b = (1+\rho)/(1-\rho)$ ; the hyperbolic forms follow from  $\tanh \tau = \sqrt{\rho}$ ). Solving  $x + \rho/x = (1+\rho)b/C = b(1-\rho)$  gives (12), and substituting  $m = b/C$  into (10) gives (13).  $\square$

*Theorem 3* (Symmetric phase diagram). For  $d = b \geq 0$  the best interior point strictly beats every alternative — exterior placements and observed anchors alike — precisely for

$$b_-(\rho) < b < e^{2\Theta}, \quad b_-(\rho) = \frac{\sqrt{\rho} \left( 2 - \sqrt{2(1-\rho)} \right)}{1+\rho}, \quad (14)$$

the lower edge being the smaller root of  $b^2(1+\rho) - 4b\sqrt{\rho} + 2\rho = 0$ . For  $b \geq e^{2\Theta}$  the optimal final location is an *observed* one: staying at an anchor beats every unvisited point, interior or exterior.

*Proof.* If  $b \geq e^{2\Theta}$ , Theorem 2 gives interior supremum  $b$ , unattained, while every exterior point loses to the anchor pointwise: for  $b > 1$  and any  $\lambda \in [0, 1)$ ,

$$b - [b\lambda + \tfrac{1}{2}(1 - \lambda^2)] = (1 - \lambda)[b - \tfrac{1+\lambda}{2}] > 0 \quad (15)$$

(the suprema coincide at  $b$ ; the dominance is strict at every unvisited point): staying at an anchor wins. If  $\sinh 2\tau < b < e^{2\Theta}$  the interior optimum is  $U_{\text{off}}$ , and both comparisons are one line: for  $b > 1$ ,

$$U_{\text{off}} - b = \tfrac{1}{2}e^{-2\Theta}(b - e^{2\Theta})^2 > 0, \quad (16)$$

the arithmetic–geometric mean inequality applied to the two terms of (13), with equality only at the tangency  $b = e^{2\Theta}$  defining the upper edge; for  $b \leq 1$ ,

$$U_{\text{off}} - \tfrac{1}{2}(b^2 + 1) = \tfrac{1}{2}(e^{2\Theta} - 1)(1 - b^2e^{-2\Theta}) > 0. \quad (17)$$

If  $b \leq \sinh 2\tau$  the optimum is the middle, with value  $U_{\text{mid}} = b \frac{2\sqrt{\rho}}{1+\rho} + \frac{1-\rho}{2(1+\rho)}$ ; for  $b \leq 1$ ,  $U_{\text{mid}} > \zeta(b)$  is the stated quadratic, and  $b_- \leq 2\sqrt{\rho}/(1+\rho) < \sinh 2\tau$  places the lower edge inside this regime. No gap opens above the middle regime: when  $\sinh 2\tau \leq 1$  (i.e.  $\rho \leq 3 - 2\sqrt{2}$ ), the larger quadratic root satisfies  $b_+ \geq \sinh 2\tau$ , equivalent to  $(1 - \rho)^3 \geq 8\rho^2$ , which holds on that range (left side decreasing, right increasing, valid at the endpoint); when  $\sinh 2\tau > 1$ ,  $b_+ \geq 1$  reduces, with  $u = \sqrt{\rho}$ , to  $2u^2(1+u) \geq (1-u)^3$ , true at  $u = \sqrt{2}-1$  (it reads  $20\sqrt{2} \geq 28$ ) and increasingly so; and on  $1 < b \leq \sinh 2\tau$ ,  $U_{\text{mid}} \geq b$  reduces to  $b \leq (1+u)/(2(1-u))$ , which covers the stretch because  $\sinh 2\tau \leq (1+u)/(2(1-u))$  is  $(1-u)^2 \geq 0$ .  $\square$

Along the two-shot-inherited bracket  $\rho = b$  the region opens at  $b^* = 0.4233\dots$ , the root of  $U_{\text{mid}}(b, b) = \zeta(b)$ . The geometry of the upper edge is easy to misread: for  $b > 1$ , interior dominance requires  $\rho > (b-1)/(b+1)$ , so higher anchors demand a larger bracket correlation, which is a *shorter* physical bracket. Very good news is forgiven only by narrow brackets; on a wide bracket, very good news means stay.

## 8 Beyond equal anchors

For general  $(b, d)$  the interior stationary points solve the quartic

$$x^4 - (b - \rho d)x^3 + \rho(d - \rho b)x - \rho^2 = 0, \quad (18)$$

obtained by differentiating (5); its real roots in  $(\rho, 1)$ , compared against  $\zeta(b)$  and  $\zeta(d)$ , decide the move exactly and replace any grid. The symmetric window does not extend indefinitely in  $d$ . At  $b = \frac{1}{2}$ ,  $\rho = \frac{1}{2}$ , the interior wins precisely for  $d \in (\approx 0.4715, 1)$ : below, the ray past the better end wins; at  $d = 1$  every strictly interior point loses, by the identity

$$L(x) - 1 = -\frac{(2x-1)^2(x^2+x+1)}{6x^2} < 0, \quad \tfrac{1}{2} < x < 1, \quad (19)$$

and for  $d \geq 1$  staying at the second anchor is optimal, since  $d$ 's coefficient in the bridge mean is below one and raising  $d$  only widens the loss. Symmetric dominance is a window, not a half-line.




Figure 2: Left: the exact symmetric revisit region (14) in the plane of bracket correlation  $\rho$  and anchor value  $b = d$ , bounded below by the quadratic root  $b_-$  and above by  $e^{2\Theta}$ . The dotted curve is the pitchfork  $b = \sinh 2\tau$ : below it the optimal interior point is the bracket middle, above it the mirror pair (12) hugging the ends. The dashed line is the two-shot bracket  $\rho = b$ , entering the region at  $b^* = 0.4233$ . Right: the value of the interior option — the percentage increase in expected payoff of the full three-shot optimum over the best policy denied interior placements (it may still place on either exterior ray) — as a function of the first discovered value  $b$ , with the bracket optimized in both cases.

## 9 The three-shot policy, computed

Writing  $V(b, d, \rho)$  for the best of  $\zeta(b)$ ,  $\zeta(d)$  and the interior optimum (via (18)), the three-shot value of a bracket  $\rho$  is  $\log \mathbb{E}_d[e^{V(b, d, \rho)}]$  with  $d \sim N(b\rho, 1 - \rho^2)$ , computed by adaptive quadrature and optimized over  $\rho$ . The comparison policy denies interior placement but keeps both exterior rays. It is not a ban on reversing direction.

Three observations, all in Table 1 and Figure 2. First, a below-median start still means flight, and the interior option is then worthless. Second, the three-shot bracket is wider than the two-shot step, but most of the widening is *not* bought by the interior option: at  $b = 0.5$  the two-shot policy steps to  $\rho = 0.5$ , excluding interior placement already widens the optimum to  $\rho = 0.323$ , and admitting it widens further to  $\rho = 0.289$ . The extra evaluation and the choice it enables account for most of the stretch; the interior option adds a further, smaller pull. Third, the interior option's worth peaks near 1.1% of expected value around  $b \approx 0.83$  and vanishes at both extremes. Excluding interior placement is thus cheap by value; it is the prescription that errs, and for equal anchors it errs exactly on the revisit window (14) — high, but below the ceiling  $e^{2\Theta}$  above which both policies stay. For unequal observations the erring region is a joint condition on  $(b, d, \rho)$ , decided by the quartic of Section 6, not membership of each value separately: at  $\rho = \frac{1}{2}$ ,  $b = \frac{1}{2}$ ,  $d = 2$ , both values lie in the symmetric window, yet staying at the second anchor is optimal and the exclusion costs nothing.



| $b$  | optimal $\rho$ | value  | value, interior excluded | interior worth |
|------|----------------|--------|--------------------------|----------------|
| 0.1  | 0.053          | 0.8396 | 0.8393                   | 0.03%          |
| 0.3  | 0.167          | 0.8723 | 0.8692                   | 0.31%          |
| 0.5  | 0.289          | 0.9386 | 0.9314                   | 0.73%          |
| 0.7  | 0.402          | 1.0375 | 1.0271                   | 1.05%          |
| 0.83 | 0.469          | 1.1183 | 1.1072                   | 1.12%          |
| 0.9  | 0.502          | 1.1670 | 1.1562                   | 1.09%          |
| 1.2  | 0.622          | 1.4059 | 1.3991                   | 0.68%          |

Table 1: The three-shot game: optimal bracket correlation, log-value, the log-value with interior placement excluded (exterior rays in both directions permitted), and the percentage the interior option adds to expected payoff. Adaptive quadrature over the second value  $d$ ; interior optima exact via the stationarity quartic (18). For  $b < 0$  the optimal bracket correlation tends to zero and the limiting log-value is exactly  $\frac{1}{2} + \log(1 + 1/\sqrt{2\pi}) = 0.8357\dots$

## 10 Discussion

Two distinct comparisons should not be merged. Against a *visited* point of matching conditional mean, any unvisited location wins by Jensen alone: it carries variance, the known value carries none. Between two *unvisited* rivals of matching mean, an interior point and its constructed exterior counterpart, Theorem 1 gives the constant variance ratio  $e^{2\Theta}$ . Neither comparison guarantees that a matched-mean unvisited point exists when the observed value is high. Mean reversion caps the conditional mean of every unvisited point. Above  $b = e^{2\Theta}$  the cap binds and the optimal final location is one already seen. Below it, the optimal location is past the better end when the news is bad, and between the anchors when the news is good.

With one evaluation remaining there is nothing to explore, since no discovery at the final point can inform a later choice. The last move is pure placement. The preference for uncertain locations comes from Jensen’s inequality, not from any value of information. Exploration proper enters only with more shots remaining. The natural conjecture from this endgame is a two-phase  $k$ -shot policy: step out past the running best while values climb, then refine between the two best once a good region is bracketed, with the pitchfork inherited by the refinement phase. There is no irreversible switch, since an adverse interior observation can make leaving attractive again. The bridge formulas make each policy evaluation cheap; the search over adaptive policies that a  $k$ -shot dynamic program requires remains to be analyzed, and we make no claim that it is small. The same placement problem arises as the line-search step inside derivative-free optimizers, where one-dimensional slices of real objectives can be rough; we pursue that use separately in the humpday benchmark package.

The problem solved here is terminal placement: the payoff is the value at the final location, not the running maximum of the path. Section 2 locates it in the literature; the short version is that the tradition’s exact objects are thresholds in  $\sigma^2/|\mu|$  on drifting Brownian paths with endogenous evaluations and flow payoffs, and the fixed-budget, terminal-payoff, mean-reverting combination, with its phase diagram, constant amplification  $e^{2\Theta}$ , and constrained parabola (10), is new.

## Reproducibility

All numbers and the figure are produced by `numerics2.py` alongside this source; the review-driven identities (the constant variance ratio, the  $d = 1$  factorization, the stationarity quartic, the exact flee limit) are independently checked in `verify/check_review.py`, and the earlier counterexample machinery lives in `verify/check_inside.py` and `verify/check_middle.py`, all in the [browniansearch](#) repository.

## References

- [1] S. Callander. Searching for good policies. *American Political Science Review*, 105(4):643–662, 2011.
- [2] S. Callander. Searching and learning by trial and error. *American Economic Review*, 101(6):2277–2308, 2011.
- [3] A. Bardhi and S. Callander. Learning in a correlated world. *Annual Review of Economics*, forthcoming.
- [4] A. Bardhi. Attributes: selective learning and influence. *Econometrica*, 92(2):311–353, 2024.
- [5] C. Urgun and L. Yariv. Contiguous search: exploration and ambition on uncharted terrain. *Journal of Political Economy*, 133, 2025.
- [6] Y. F. Wong. Forward-looking experimentation of correlated alternatives. *Theoretical Economics*, 20, 2025.
- [7] S. Callander. A theory of policy expertise. *Quarterly Journal of Political Science*, 3(2):123–140, 2008.
- [8] S. Callander and P. Hummel. Preemptive policy experimentation. *Econometrica*, 82(4):1509–1528, 2014.
- [9] S. Callander and T. S. Clark. Precedent and doctrine in a complicated world. *American Political Science Review*, 111(1):184–203, 2017.
- [10] S. Callander and N. Matouschek. The risk of failure: trial and error learning and long-run performance. *American Economic Journal: Microeconomics*, 11(1):44–78, 2019.
- [11] S. Callander and N. McCarty. Agenda control under policy uncertainty. *American Journal of Political Science*, 68(1):210–226, 2024.
- [12] S. Callander, N. S. Lambert, and N. Matouschek. Innovation and competition on a rugged technological landscape. *American Economic Journal: Microeconomics*, forthcoming.
- [13] Y. C. Aybas and S. Callander. Cheap talk in complex environments. Working paper, 2025.
- [14] S. C. Ganz. Hyperopic search: organizations learning about managers learning about strategies. *Organization Science*, 31(4):821–838, 2020.
- [15] U. Garfagnini and B. Strulovici. Social experimentation with interdependent and expanding technologies. *Review of Economic Studies*, 83(4):1579–1613, 2016.

- [16] D. Cetemen, C. Urgun, and L. Yariv. Collective progress: dynamics of exit waves. *Journal of Political Economy*, 131(9):2402–2450, 2023.
- [17] C. Carnehl and J. Schneider. A quest for knowledge. *Econometrica*, 93(2):623–659, 2025.
- [18] C. Hodgson and G. Lewis. You can lead a horse to water: spatial learning and path dependence in consumer search. *Econometrica*, 93(4):1299–1332, 2025.
- [19] D. Glick and C. D. Myers. Learning from others: an experimental test of Brownian motion uncertainty models. *Journal of Theoretical Politics*, 27(4):588–612, 2015.
- [20] S. Malladi. Searching in the dark and learning where to look. Working paper, 2025.
- [21] M. Banchio and S. Malladi. Rediscovery. Working paper, arXiv:2504.19761, 2025.
- [22] J. Grosse, C. Zhang, and P. Hennig. Optimistic optimization of Gaussian process samples. *Transactions on Machine Learning Research*, 2023.
- [23] M. A. Butler and A. A. King. Phylogenetic comparative analysis: a modeling approach for adaptive evolution. *The American Naturalist*, 164(6):683–695, 2004.
- [24] T. F. Hansen. Stabilizing selection and the comparative analysis of adaptation. *Evolution*, 51(5):1341–1351, 1997.
- [25] C. E. Rasmussen and C. K. I. Williams. *Gaussian Processes for Machine Learning*. MIT Press, 2006.
- [26] J. Snoek, H. Larochelle, and R. P. Adams. Practical Bayesian optimization of machine learning algorithms. In *Advances in Neural Information Processing Systems 25*, 2012.
- [27] E. Weinberger. Correlated and uncorrelated fitness landscapes and how to tell the difference. *Biological Cybernetics*, 63(5):325–336, 1990.
- [28] X. Feng, K. A. Nguyen, and Z. Luo. WiFi RTT RSS dataset for indoor positioning. Zenodo, 2024. doi:10.5281/zenodo.11558192.
- [29] M. Gudmundson. Correlation model for shadow fading in mobile radio systems. *Electronics Letters*, 27(23):2145–2146, 1991.
- [30] E. Ries. *The Lean Startup*. Crown Business, New York, 2011.
- [31] A. Camuffo, A. Cordova, A. Gambardella, and C. Spina. A scientific approach to entrepreneurial decision making: evidence from a randomized control trial. *Management Science*, 66(2):564–586, 2020.
- [32] A. Camuffo, A. Gambardella, D. Messinese, E. Novelli, E. Paolucci, and C. Spina. A scientific approach to entrepreneurial decision making: large scale replication and extension. *Strategic Management Journal*, 45(6):1209–1237, 2024.
- [33] O. Baumann, J. Schmidt, and N. Stieglitz. Effective search in rugged performance landscapes: a review and outlook. *Journal of Management*, 45(1):285–318, 2019.

- [34] J.-B. Grill, M. Valko, and R. Munos. Optimistic optimization of a Brownian. In *Advances in Neural Information Processing Systems 31*, 2018.
- [35] J. M. Calvin. Average performance of a class of adaptive algorithms for global optimization. *Annals of Applied Probability*, 7(3):711–730, 1997.
- [36] A. Alabert and R. Caballero. On the minimum of a conditioned Brownian bridge. *Stochastic Models*, 34(3):269–291, 2018.
- [37] R. Agrawal. The continuum-armed bandit problem. *SIAM Journal on Control and Optimization*, 33(6):1926–1951, 1995.